

Infron vs OpenRouter Routing, Cache, Cost, and Streaming Performance A/B Test Report

Abstract and Executive Outline

Keywords: Prompt Caching; A/B Testing; Provider Routing; Cache Affinity; Latency; Throughput; Cost Attribution; kimi-k2.5

Abstract

This report evaluates `moonshotai/kimi-k2.5` on Infron and OpenRouter across cache reuse, observed cost, throughput, E2E latency, and Streaming TTFT under Prompt Caching workloads.

The main findings are: Infron leads Cache hit rate in 0/4 routing modes; OpenRouter leads in 2/4; 2/4 tie; Infron leads Observed cost in 1/4 routing modes; OpenRouter leads in 1/4; 2/4 tie; Infron leads Throughput in 0/4 routing modes; OpenRouter leads in 2/4; 2/4 tie; Infron leads E2E latency in 0/4 routing modes; OpenRouter leads in 2/4; 2/4 tie; Infron leads Streaming TTFT in 0/4 routing modes; OpenRouter leads in 2/4; 2/4 tie.

Overall, Infron does not show a clear cross-mode strength in this run, while OpenRouter's cross-mode strengths are throughput, Streaming TTFT, E2E latency, cache reuse. Platform choice should be driven by workload objective rather than a single headline metric.

Figure 0: Normalized Capability Radar

The five radar axes represent throughput, price, E2E latency, Streaming TTFT, and cache hit rate. Every metric is normalized to a 0–100 score, with farther outward meaning better.

Thick solid lines show platform–level contours; translucent lines and points show individual routing modes.

Conclusion Overview: Core Metric Winners by Routing Mode

Based on strict A/B paired samples. Blue represents Infron, orange represents OpenRouter; gold cells mark the winner for the routing objective.

Throughput	Observed cost	E2E latency	Streaming TTFT	Cache hit rate
Tie in 4/4 modes Max advantage 4.30%, higher is better	Tie in 4/4 modes Max advantage 5.10%, lower is better	Tie in 4/4 modes Max advantage 105.13%, lower is better	Tie in 4/4 modes Max advantage 135.09%, lower is better	Tie in 4/4 modes Max advantage 5.46%, higher is better

Routing mode	Throughput objective	Cost objective	Latency objective	TTFT objective	Cache result
Throughput First throughput	OpenRouter advantage 2.51%	OpenRouter advantage 5.10%	OpenRouter advantage 8.10%	OpenRouter advantage 48.82%	OpenRouter advantage 5.46%
Price First price	Tie tie	Tie tie	Tie tie	Tie tie	Tie tie
Latency First latency	Tie tie	Tie tie	Tie tie	Tie tie	Tie tie
TTFT First ttft	OpenRouter advantage 4.30%	Infron advantage 73.04%	OpenRouter advantage 105.13%	OpenRouter advantage 135.09%	OpenRouter advantage 0.77%

Executive Outline

Dimension	Conclusion	Evidence
Controls	First/second <code>usage.prompt_tokens</code> deltas are limited to 50 tokens within each <code>sort/group/round</code> pair.	Methods and data quality
Cache reuse	Infron leads Cache hit rate in 0/4 routing modes; OpenRouter leads in 2/4; 2/4 tie	Overall metrics and mechanism section
Observed cost	Infron leads Observed cost in 1/4 routing modes; OpenRouter leads in 1/4; 2/4 tie	Overall metrics and provider drill-down
Performance	Infron leads Throughput in 0/4 routing modes; OpenRouter leads in 2/4; 2/4 tie; Infron leads E2E latency in 0/4 routing modes; OpenRouter leads in 2/4; 2/4 tie; Infron leads Streaming TTFT in 0/4 routing modes; OpenRouter leads in 2/4; 2/4 tie	Charts and statistical tests
Attribution boundary	Claims use observable response telemetry: usage, cost, TTFT, latency, provider fields, and cache tokens.	Provider/Route drill-down
Business meaning	Long-context, RAG-prefix, agent-tool, and high-frequency template workloads should evaluate cache rate, cost, first-token latency, and E2E latency together.	Discussion and conclusion

Routing–Mode Conclusions

Routing mode	Objective	Cache winner	Cost winner	Throughput winner	E2E latency winner	TTFT winner	Interpretation
Throughput First	Maximize output capacity per unit time	OpenRouter	OpenRouter	OpenRouter	OpenRouter	OpenRouter	OpenRouter leads overall (5/5 metrics)
Price First	Minimize request and token cost	Tie	Tie	Tie	Tie	Tie	Mixed result
Latency First	Minimize full–response waiting time	Tie	Tie	Tie	Tie	Tie	Mixed result
TTFT First	Minimize streaming first–token time	OpenRouter	Infron	OpenRouter	OpenRouter	OpenRouter	OpenRouter leads overall (4/5 metrics)

1. Introduction: Background, Questions, and Contributions

LLM inference–platform behavior is shaped by provider routing, prompt caching, streaming response handling, cost attribution, and fallback policy. This report treats the platform as an observable system and measures speed, cost, cache reuse, and first–token experience through strict A/B pairs.

1.1 Research Hypotheses

Hypothesis	Statement	Validation metric
H1	Stronger provider/cache affinity improves token–level cache hit rate for repeated stable prefixes.	Second–request cache–read tokens and token cache hit rate
H2	Higher cache hit rate can reduce observed cost, but does not necessarily reduce TTFT or E2E latency.	Observed cost, average TTFT, average request latency
H3	Different routing sorts change provider selection and produce different cost/throughput/latency frontiers.	Provider distribution, throughput, latency, cost

1.2 Contributions

- Uses response–returned `usage.prompt_tokens` as the input–token control while allowing small cross–platform accounting variance up to 50 tokens.
- Extends prompt–caching evaluation to cost, throughput, E2E latency, TTFT, provider distribution, reasoning telemetry, and paired statistical tests.
- Scopes every claim to observable response telemetry instead of private routing internals.


2. Experimental Design, Dataset, and Controls

2.1 Dataset Construction

The dataset is `controlled_cache_probe`, covering 4 routing sorts, 2 platforms, 4 groups, and 50 rounds per group. Each round sends identical first/second Chat Completions requests; the second request observes cache-read tokens, TTFT, and E2E latency.

The built-in business templates represent stable long-context workloads such as RAG support, agent tool instructions, marketing automation, and code review.

2.2 Controlled Variables



Controlled-variable rule: within the same `sort/group/round`, both platforms must have first/second `usage.prompt_tokens` deltas no greater than 50 tokens. Total Input Tokens are computed from response-returned usage, not local tokenizer estimates.

2.3 Metric Definitions

Metric	Definition	Direction
Call cache hit rate	Share of second requests with <code>cache_read_tokens > 0</code>	Higher is better
Token cache hit rate	Second-request cache-read tokens / second-request prompt tokens	Higher is better
Observed cost	Sum of first + second request usage/cost values	Lower is better
Throughput	Completion tokens / request latency seconds	Higher is better
E2E latency	Full response latency per request	Lower is better
TTFT	Streaming first-token arrival time	Lower is better
Reasoning telemetry	Reasoning tokens from response usage	Used to explain latency, throughput, and cost

3. Experimental Environment and Data Quality

Item	Configuration
Model	moonshotai/kimi-k2.5
Provider model IDs	infron: moonshotai/kimi-k2.5 ; openrouter: moonshotai/kimi-k2.5
Platforms	Infron and OpenRouter
API protocol	/v1/chat/completions
Routing modes	Throughput First, Price First, Latency First, TTFT First
Groups	4
Rounds per group	50
Workers	24
Request mode	Streaming Chat Completions with TTFT collection
Reasoning / thinking control	No explicit reasoning/thinking parameter; model and platform defaults are preserved
Prompt length tiers	short ≈1500, medium ≈8000, long ≈32000
Excluded records	1450

4. Results: Overall Metrics and Main Findings

Routing mode	Platform	Strict pairs	Total Input Tokens	Token cache hit rate	Observed cost	Throughput	E2E latency	Streaming TTFT
Throughput First	Infron	2	21804	15.26%	\$0.01313000	1.13 tok/s	3529.01 ms	3362.17 ms
Throughput First	OpenRouter	2	21800	98.64%	\$0.00215256	3.98 tok/s	3264.43 ms	2259.22 ms
Price First	Infron	0	0	0.00%	\$0.00000000	0.00 tok/s	0.00 ms	0.00 ms
Price First	OpenRouter	0	0	0.00%	\$0.00000000	0.00 tok/s	0.00 ms	0.00 ms
Latency First	Infron	0	0	0.00%	\$0.00000000	0.00 tok/s	0.00 ms	0.00 ms
Latency First	OpenRouter	0	0	0.00%	\$0.00000000	0.00 tok/s	0.00 ms	0.00 ms
TTFT First	Infron	73	2203077	98.84%	\$0.14219200	0.89 tok/s	6604.17 ms	6317.80 ms
TTFT First	OpenRouter	73	2202954	99.60%	\$0.24604531	4.72 tok/s	3219.54 ms	2687.39 ms

4.1 Tail Latency and Statistical Tests

Tail percentiles expose risk that averages hide.

Routing mode	Platform	P50 latency	P95 latency	P99 latency	P50 TTFT	P95 TTFT	P99 TTFT
Throughput First	Infron	3575.51 ms	4584.96 ms	4631.99 ms	3342.05 ms	4476.37 ms	4525.71 ms
Throughput First	OpenRouter	3106.35 ms	5481.42 ms	5622.76 ms	2229.86 ms	3348.97 ms	3425.96 ms
Price First	Infron	N/A	N/A	N/A	N/A	N/A	N/A
Price First	OpenRouter	N/A	N/A	N/A	N/A	N/A	N/A
Latency First	Infron	N/A	N/A	N/A	N/A	N/A	N/A
Latency First	OpenRouter	N/A	N/A	N/A	N/A	N/A	N/A
TTFT First	Infron	4544.57 ms	15517.55 ms	20824.11 ms	4335.68 ms	15382.81 ms	19942.85 ms
TTFT First	OpenRouter	2527.21 ms	8827.67 ms	9735.58 ms	2061.95 ms	8105.85 ms	9149.96 ms

Mean deltas use bootstrap 95% CIs; p-values use paired sign-flip permutation tests.

Routing mode	Metric	Mean delta	95% CI	p-value	Pairs	Interpretation
Throughput First	Latency: OpenRouter – Infron	–529.16 ms	–1505.05 ms to 446.73 ms	1.0000	2	Positive means lower Infron latency
Throughput First	TTFT: OpenRouter – Infron	–2205.89 ms	–4205.20 ms to –206.57 ms	0.4994	2	Positive means lower Infron TTFT
Throughput First	Throughput: Infron – OpenRouter	–2.8370 tok/s	–3.0554 tok/s to –2.6187 tok/s	0.4994	2	Positive means higher Infron throughput
Throughput First	Cost: OpenRouter – Infron	\$–0.00548872	\$–0.00926436 to \$–0.00171308	0.4994	2	Positive means lower Infron cost
Throughput First	Token Cache Hit: Infron – OpenRouter	–49.76 pp	–99.47 pp to –0.05 pp	0.4994	2	Positive means higher Infron cache hit
Price First	Latency: OpenRouter – Infron	N/A	N/A to N/A	N/A	0	Positive means lower Infron latency
Price First	TTFT: OpenRouter – Infron	N/A	N/A to N/A	N/A	0	Positive means lower Infron TTFT
Price First	Throughput: Infron – OpenRouter	N/A	N/A to N/A	N/A	0	Positive means higher Infron throughput
Price First	Cost: OpenRouter – Infron	N/A	N/A to N/A	N/A	0	Positive means lower Infron cost
Price First	Token Cache Hit: Infron – OpenRouter	N/A	N/A to N/A	N/A	0	Positive means higher Infron cache hit
Latency First	Latency: OpenRouter – Infron	N/A	N/A to N/A	N/A	0	Positive means lower Infron latency
Latency First	TTFT: OpenRouter – Infron	N/A	N/A to N/A	N/A	0	Positive means lower Infron TTFT
Latency First	Throughput: Infron – OpenRouter	N/A	N/A to N/A	N/A	0	Positive means higher Infron throughput

Routing mode	Metric	Mean delta	95% CI	p-value	Pairs	Interpretation
Latency First	Cost: OpenRouter – Infron	N/A	N/A to N/A	N/A	0	Positive means lower Infron cost
Latency First	Token Cache Hit: Infron – OpenRouter	N/A	N/A to N/A	N/A	0	Positive means higher Infron cache hit
TTFT First	Latency: OpenRouter – Infron	–6769.25 ms	–8636.13 ms to –5068.44 ms	<0.001	73	Positive means lower Infron latency
TTFT First	TTFT: OpenRouter – Infron	–7260.84 ms	–9115.20 ms to –5591.96 ms	<0.001	73	Positive means lower Infron TTFT
TTFT First	Throughput: Infron – OpenRouter	–4.2878 tok/s	–5.0085 tok/s to –3.6878 tok/s	<0.001	73	Positive means higher Infron throughput
TTFT First	Cost: OpenRouter – Infron	\$0.00142265	\$0.00096758 to \$0.00194698	<0.001	73	Positive means lower Infron cost
TTFT First	Token Cache Hit: Infron – OpenRouter	–1.50 pp	–4.43 pp to 0.13 pp	0.2369	73	Positive means higher Infron cache hit

4.2 Reasoning / Thinking Control Check

This run does not explicitly set reasoning/thinking parameters and keeps model/platform defaults; this table records reasoning telemetry under default behavior.

Routing mode	Platform	Reasoning tokens	Avg reasoning tokens/request	Reasoning requests	Avg first reasoning token	Avg TTFT	Avg E2E latency
Throughput First	Infron	0	0.0000	0	0.00 ms	3362.17 ms	3529.01 ms
Throughput First	OpenRouter	52	13.0000	4	2259.22 ms	2259.22 ms	3264.43 ms
Price First	Infron	0	0.0000	0	0.00 ms	0.00 ms	0.00 ms
Price First	OpenRouter	0	0.0000	0	0.00 ms	0.00 ms	0.00 ms
Latency First	Infron	0	0.0000	0	0.00 ms	0.00 ms	0.00 ms
Latency First	OpenRouter	0	0.0000	0	0.00 ms	0.00 ms	0.00 ms
TTFT First	Infron	0	0.0000	0	8054.65 ms	6317.80 ms	6604.17 ms
TTFT First	OpenRouter	2207	15.1164	146	2687.39 ms	2687.39 ms	3219.54 ms

4.3 API Protocol Compatibility Matrix

This run uses `/v1/chat/completions`; this table records success response, usage, cost, and cache telemetry coverage for both platforms under that protocol.

API protocol	Endpoint	Platform	Pairs	Requests	Success rate	Usage coverage	Token usage coverage	Cost coverage	Cache telemetry coverage
chat_completions	/v1/chat/completions	Infron	800	1600	33.62%	100.00%	100.00%	100.00%	100.00%
chat_completions	/v1/chat/completions	OpenRouter	800	1600	30.50%	100.00%	100.00%	100.00%	100.00%

5. Result Visualizations by Routing Mode

The following charts are driven by the strict-paired summary data and present routing-mode, platform, cost, cache, E2E latency, Streaming TTFT, and upstream provider differences.

5.1 Core Metric Chart Overview

The charts below are generated from the same strict-paired summary data.

**Figure 3–E2:
Normalized
capability
contour**

Five metrics are normalized to 0–100, where higher is better.

**Figure 3–E3:
Cost–cache
efficiency plane**

X-axis is token cache hit rate; Y-axis is observed total cost.

**Figure
3–
E1:
Routing–
mode
core
metric
comparison**

Infron and OpenRouter are shown side by side by routing mode.

Figure 3–E4: E2E latency and Streaming TTFT

Solid lines show E2E latency; dashed lines show Streaming TTFT.

Figure 3–E5: Upstream provider distribution

Provider shares explain cache–domain and performance differences.

Routing–Mode Drill–Down Charts

The horizontal axis shows relative advantage: blue to the right indicates Infron advantage, orange to the left indicates OpenRouter advantage.

Figure 4–E1: Throughput First core metric comparison

Direction and magnitude of relative advantage across five metrics.

Figure 4–E2: Price First core metric comparison

Direction and magnitude of relative advantage across five metrics.


Figure 4–E3: Latency First core metric comparison

Direction and magnitude of relative advantage across five metrics.

Figure 4–E4: TTFT First core metric comparison

Direction and magnitude of relative advantage across five metrics.

6. Infron Technical Architecture and Cache/Cost Mechanism



6.1 Multi-Provider Routing and Observable Control Plane

Requests enter a unified API, and the routing layer selects healthy upstream paths according to throughput, price, latency, or ttft objectives. Whether stable long prefixes stay in the same cache domain directly affects second-request cache-read tokens.

6.2 Provider Stick and Cache Affinity

Provider stick is cache affinity, not a permanent provider lock. Its goal is to reduce cache-domain fragmentation while respecting health and SLA constraints.

6.3 Cost-Control Path

Mechanism	Cache-rate effect	Cost effect	Observable signal
Stable prefix detection	Repeated prefixes are more likely to hit cache	Reduces repeated prefill cost	Payload SHA-256 and second-request cache-read tokens
Provider stick / cache affinity	Reduces cross-domain cache fragmentation	Avoids repeated cache warm-up	Provider distribution and token cache hit rate
Health checks and fallback	Protects availability while sometimes sacrificing cache	Reduces failure cost	HTTP status, provider distribution, tail latency
Cost-aware routing	Prefers lower-cost paths under constraints	Reduces total and per-round cost	Observed cost, cost breakdown coverage, cache-read tokens

7. Provider/Route Drill-Down

Routing mode	Platform	Total requests	Attributed requests	Provider distribution
Throughput First	Infron	4	4	novita 4

Routing mode	Platform	Total requests	Attributed requests	Provider distribution
Throughput First	OpenRouter	4	4	ModelRun 4
Price First	Infron	0	0	Not returned
Price First	OpenRouter	0	0	Not returned
Latency First	Infron	0	0	Not returned
Latency First	OpenRouter	0	0	Not returned
TTFT First	Infron	146	146	alibaba/cn 123, deepinfra 23
TTFT First	OpenRouter	146	146	ModelRun 116, DigitalOcean 26, DeepInfra 2, SiliconFlow 2

Upstream Provider Detail Distribution

Routing mode	Platform	Upstream provider	Requests	Share	first/ second	Covered rounds	Avg TTFT	Avg latency	Prompt tokens	Completion tokens
Throughput First	Infron	novita	4	100.00%	2/2	2	3362.17 ms	3529.01 ms	21804	16
Throughput First	OpenRouter	ModelRun	4	100.00%	2/2	2	2259.22 ms	3264.43 ms	21800	52
Price First	Infron	Not returned	0	0.00%	0/0	0	N/A	N/A	0	0
Price First	OpenRouter	Not returned	0	0.00%	0/0	0	N/A	N/A	0	0
Latency First	Infron	Not returned	0	0.00%	0/0	0	N/A	N/A	0	0
Latency First	OpenRouter	Not returned	0	0.00%	0/0	0	N/A	N/A	0	0
TTFT First	Infron	alibaba/cn	123	84.25%	65/58	68	6037.20 ms	6199.51 ms	1888563	492
TTFT First	Infron	deepinfra	23	15.75%	8/15	18	7818.43 ms	8768.19 ms	314514	368
TTFT First	OpenRouter	ModelRun	116	79.45%	61/55	61	2632.83 ms	2986.61 ms	1757551	1514
TTFT First	OpenRouter	DigitalOcean	26	17.81%	10/16	16	2296.73 ms	3405.60 ms	369149	416
TTFT First	OpenRouter	DeepInfra	2	1.37%	1/1	1	11668.90 ms	11906.07 ms	72726	32
TTFT First	OpenRouter	SiliconFlow	2	1.37%	1/1	1	1948.52 ms	5624.29 ms	3528	259

7.1 Cache–Rate and Cost Divergence Drill–Down

This table combines cache, cost, provider distribution, and reasoning telemetry to explain routing–mode differences.

Routing mode	Cache–hit delta	Infron cost multiple	Infron top path	OpenRouter top path	Reasoning token delta	Main attribution
Throughput First	–83.38 pp	6.10x	novita 100.00%	ModelRun 100.00%	–52	OpenRouter has higher cache and lower cost; provider/cache–domain mix is the main signal
Price First	+0.00 pp	N/A	未返回	未返回	+0	Cache and cost move in different directions; evaluate with speed metrics
Latency First	+0.00 pp	N/A	未返回	未返回	+0	Cache and cost move in different directions; evaluate with speed metrics
TTFT First	–0.76 pp	0.58x	alibaba/cn 84.25%	ModelRun 79.45%	–2207	Cache and cost move in different directions; evaluate with speed metrics

8. Stratified Results: Prompt–Length Cache Performance

This section aggregates second–request cache–read tokens, token–level cache hit rate, observed cost, E2E latency, and Streaming TTFT by prompt–length tier. Bold cells mark the advantaged side within each tier.

Prompt–Length Tier Overview

Prompt length tier	Target tokens	Platform	Pairs	Second prompt tokens	Second cache read tokens	Token cache hit rate	Observed cost	Avg E2E latency	Avg TTFT
short	1500	Infron	27	47650	45312	95.09%	\$0.00995100	5294.39 ms	4885.33 ms
short	1500	OpenRouter	27	47628	45504	95.54%	\$0.01343926	2545.05 ms	1971.02 ms
medium	8000	Infron	25	228419	209184	91.58%	\$0.04532000	6429.55 ms	6203.25 ms
medium	8000	OpenRouter	25	228400	227392	99.56%	\$0.05440489	3054.37 ms	2516.64 ms
long	32000	Infron	23	836368	835872	99.94%	\$0.10005100	8064.13 ms	7866.91 ms
long	32000	OpenRouter	23	836349	834928	99.83%	\$0.18035371	4194.78 ms	3676.69 ms

Prompt Length x Routing Mode Cache Hit Rate

Prompt length tier	Routing mode	Infron	OpenRouter	Winner
short	Throughput First	94.28%	94.33%	OpenRouter
short	Price First	0.00%	0.00%	tie
short	Latency First	0.00%	0.00%	tie
short	TTFT First	95.12%	95.59%	OpenRouter
medium	Throughput First	0.00%	99.47%	OpenRouter
medium	Price First	0.00%	0.00%	tie
medium	Latency First	0.00%	0.00%	tie
medium	TTFT First	95.39%	99.56%	OpenRouter
long	Throughput First	0.00%	0.00%	tie
long	Price First	0.00%	0.00%	tie
long	Latency First	0.00%	0.00%	tie
long	TTFT First	99.94%	99.83%	Infron

9. Stratified Results: Group–Level Stability

Throughput First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	0	0	0.00%	\$0.00000000	N/A	N/A
Infron	2	2	2	15.26%	\$0.01313000	4584.96 ms	4476.37 ms
Infron	3	0	0	0.00%	\$0.00000000	N/A	N/A
Infron	4	0	0	0.00%	\$0.00000000	N/A	N/A
OpenRouter	1	0	0	0.00%	\$0.00000000	N/A	N/A
OpenRouter	2	2	2	98.64%	\$0.00215256	5481.42 ms	3348.97 ms
OpenRouter	3	0	0	0.00%	\$0.00000000	N/A	N/A
OpenRouter	4	0	0	0.00%	\$0.00000000	N/A	N/A

Price First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	0	0	0.00%	\$0.00000000	N/A	N/A
Infron	2	0	0	0.00%	\$0.00000000	N/A	N/A
Infron	3	0	0	0.00%	\$0.00000000	N/A	N/A

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	4	0	0	0.00%	\$0.00000000	N/A	N/A
OpenRouter	1	0	0	0.00%	\$0.00000000	N/A	N/A
OpenRouter	2	0	0	0.00%	\$0.00000000	N/A	N/A
OpenRouter	3	0	0	0.00%	\$0.00000000	N/A	N/A
OpenRouter	4	0	0	0.00%	\$0.00000000	N/A	N/A

Latency First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	0	0	0.00%	\$0.00000000	N/A	N/A
Infron	2	0	0	0.00%	\$0.00000000	N/A	N/A
Infron	3	0	0	0.00%	\$0.00000000	N/A	N/A
Infron	4	0	0	0.00%	\$0.00000000	N/A	N/A
OpenRouter	1	0	0	0.00%	\$0.00000000	N/A	N/A
OpenRouter	2	0	0	0.00%	\$0.00000000	N/A	N/A
OpenRouter	3	0	0	0.00%	\$0.00000000	N/A	N/A
OpenRouter	4	0	0	0.00%	\$0.00000000	N/A	N/A

TTFT First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	3	3	99.46%	\$0.00471600	29866.65 ms	29687.25 ms
Infron	2	2	2	99.87%	\$0.00527600	4976.74 ms	4781.97 ms
Infron	3	18	18	99.57%	\$0.02929800	15322.95 ms	15190.38 ms
Infron	4	50	50	98.50%	\$0.10290200	15592.71 ms	15468.08 ms
OpenRouter	1	3	3	99.47%	\$0.01569255	3869.88 ms	3196.18 ms
OpenRouter	2	2	2	99.59%	\$0.00840456	2719.96 ms	2629.10 ms
OpenRouter	3	18	18	99.57%	\$0.04695015	5050.03 ms	4569.76 ms
OpenRouter	4	50	50	99.62%	\$0.17499804	9080.67 ms	8423.71 ms

10. Discussion: Business Value, Boundaries, and Engineering Implications

Business decisions should not rely on one metric. Stable long-context and high-frequency template workloads should prioritize cache rate and cost; realtime interaction must constrain TTFT and E2E latency; batch processing often prioritizes throughput and failure cost.

Routing mode	Business objective	Observed result	Scenarios	Caveat
Throughput First	Maximize output capacity per unit time	OpenRouter leads overall (5/5 metrics)	Batch generation, offline summaries, backend processing	First-token and E2E response are faster; check cache and cost
Price First	Minimize request and token cost	Mixed result	High-frequency templates, support automation, marketing, RAG prefixes	Decide with budget, SLA, and cache stability together
Latency First	Minimize full-response waiting time	Mixed result	Online chat, agent chains, IDE/writing assistants, realtime tools	Decide with budget, SLA, and cache stability together
TTFT First	Minimize streaming first-token time	OpenRouter leads overall (4/5 metrics)	Streaming chat, realtime copilots, first-screen feedback	First-token and E2E response are faster; check cache and cost

11. Conclusion

Routing-Mode Conclusions

Routing mode	Objective	Cache winner	Cost winner	Throughput winner	E2E latency winner	TTFT winner	Interpretation
Throughput First	Maximize output capacity per unit time	OpenRouter	OpenRouter	OpenRouter	OpenRouter	OpenRouter	OpenRouter leads overall (5/5 metrics)
Price First	Minimize request and token cost	Tie	Tie	Tie	Tie	Tie	Mixed result
Latency First	Minimize full-response waiting time	Tie	Tie	Tie	Tie	Tie	Mixed result
TTFT First	Minimize streaming first-token time	OpenRouter	Infron	OpenRouter	OpenRouter	OpenRouter	OpenRouter leads overall (4/5 metrics)

12. Limitations, Missing Data, and Next Experiments

Limitation	Impact	Next step	Current handling
Full routing trace	Cannot prove every provider choice, fallback, and retry path hop by hop	Add provider routing trace, decision logs, and fallback reasons	Use only returned provider fields and provider distribution
Longer time window	4x50 observes short-window stability but not day-level drift	Add soak tests and repeated windows	Scope conclusions to this run

Limitation	Impact	Next step	Current handling
Production corpus	Built-in templates do not cover every workload distribution	Use sanitized production-stratified corpora	Discuss representative long-context templates only
Cost-field consistency	Cost coverage and semantics may differ by platform	Reconcile with billing and provider cost breakdown	Use only explicitly returned cost fields

13. Reproducibility Appendix

Artifact	Path
Summary	export/kimi_k25_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions summary.json
Paired dataset	export/kimi_k25_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions benchmark_pairs.csv
Request-level dataset	export/kimi_k25_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions benchmark_requests.jsonl
Filtered structured records	export/kimi_k25_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions records.json
Excluded-record audit	export/kimi_k25_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions records_excluded.json
Test source	tests/test_rerun_routing_sort_cache_cost_ab.py
Benchmark runner source	scripts/rerun_routing_sort_cache_cost_ab.py
HTML report renderer source	scripts/render_glm52_deepseek_style_report.py
Dataset reference	business_representative built-in representative business templates; request-level export is benchmark_requests