

Infron vs OpenRouter Routing, Cache, Cost, and Streaming Performance A/B Test Report

Abstract and Executive Outline

Keywords: Prompt Caching; A/B Testing; Provider Routing; Cache Affinity; Latency; Throughput; Cost Attribution; kimi-k2.6

Abstract

This report evaluates `moonshotai/kimi-k2.6` on Infron and OpenRouter across cache reuse, observed cost, throughput, E2E latency, and Streaming TTFT under Prompt Caching workloads.

The main findings are: Infron leads Cache hit rate in all routing modes; Infron leads Observed cost in 3/4 routing modes; OpenRouter leads in 1/4; 0/4 tie; Infron leads Throughput in 1/4 routing modes; OpenRouter leads in 3/4; 0/4 tie; OpenRouter leads E2E latency in all routing modes; OpenRouter leads Streaming TTFT in all routing modes.

Overall, Infron's cross-mode strengths are observed cost, cache reuse, while OpenRouter's cross-mode strengths are throughput, Streaming TTFT, E2E latency. Platform choice should be driven by workload objective rather than a single headline metric.

Figure 0: Normalized Capability Radar

The five radar axes represent throughput, price, E2E latency, Streaming TTFT, and cache hit rate. Every metric is normalized to a 0–100 score, with farther outward meaning better.

Thick solid lines show platform–level contours; translucent lines and points show individual routing modes.

Conclusion Overview: Core Metric Winners by Routing Mode

Based on strict A/B paired samples. Blue represents Infron, orange represents OpenRouter; gold cells mark the winner for the routing objective.

Throughput OpenRouter wins 3/4 Max advantage 129.56%, higher is better	Observed cost Infron wins 3/4 Max advantage 67.22%, lower is better	E2E latency OpenRouter wins 4/4 Max advantage 71.95%, lower is better	Streaming TTFT OpenRouter wins 4/4 Max advantage 74.06%, lower is better	Cache hit rate Infron wins 4/4 Max advantage 13.56%, higher is better
---	--	--	---	--

Routing mode	Throughput objective	Cost objective	Latency objective	TTFT objective	Cache result
Throughput First throughput	Infron advantage 129.56%	OpenRouter advantage 32.92%	OpenRouter advantage 52.21%	OpenRouter advantage 31.71%	Infron advantage 1.23%
Price First price	OpenRouter advantage 86.92%	Infron advantage 36.21%	OpenRouter advantage 71.95%	OpenRouter advantage 74.06%	Infron advantage 11.01%
Latency First latency	OpenRouter advantage 29.53%	Infron advantage 53.16%	OpenRouter advantage 30.73%	OpenRouter advantage 27.97%	Infron advantage 13.09%
TTFT First ttft	OpenRouter advantage 51.36%	Infron advantage 67.22%	OpenRouter advantage 52.68%	OpenRouter advantage 39.66%	Infron advantage 13.56%

Executive Outline

Dimension	Conclusion	Evidence
Controls	First/second <code>usage.prompt_tokens</code> deltas are limited to 50 tokens within each <code>sort/group/round</code> pair.	Methods and data quality
Cache reuse	Infron leads Cache hit rate in all routing modes	Overall metrics and mechanism section
Observed cost	Infron leads Observed cost in 3/4 routing modes; OpenRouter leads in 1/4; 0/4 tie	Overall metrics and provider drill-down
Performance	Infron leads Throughput in 1/4 routing modes; OpenRouter leads in 3/4; 0/4 tie; OpenRouter leads E2E latency in all routing modes; OpenRouter leads Streaming TTFT in all routing modes	Charts and statistical tests
Attribution boundary	Claims use observable response telemetry: usage, cost, TTFT, latency, provider fields, and cache tokens.	Provider/Route drill-down
Business meaning	Long-context, RAG-prefix, agent-tool, and high-frequency template workloads should evaluate cache rate, cost, first-token latency, and E2E latency together.	Discussion and conclusion

Routing–Mode Conclusions

Routing mode	Objective	Cache winner	Cost winner	Throughput winner	E2E latency winner	TTFT winner	Interpretation
Throughput First	Maximize output capacity per unit time	Infron	OpenRouter	Infron	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)
Price First	Minimize request and token cost	Infron	Infron	OpenRouter	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)
Latency First	Minimize full–response waiting time	Infron	Infron	OpenRouter	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)
TTFT First	Minimize streaming first–token time	Infron	Infron	OpenRouter	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)

1. Introduction: Background, Questions, and Contributions

LLM inference–platform behavior is shaped by provider routing, prompt caching, streaming response handling, cost attribution, and fallback policy. This report treats the platform as an observable system and measures speed, cost, cache reuse, and first–token experience through strict A/B pairs.

1.1 Research Hypotheses

Hypothesis	Statement	Validation metric
H1	Stronger provider/cache affinity improves token–level cache hit rate for repeated stable prefixes.	Second–request cache–read tokens and token cache hit rate
H2	Higher cache hit rate can reduce observed cost, but does not necessarily reduce TTFT or E2E latency.	Observed cost, average TTFT, average request latency
H3	Different routing sorts change provider selection and produce different cost/throughput/latency frontiers.	Provider distribution, throughput, latency, cost

1.2 Contributions

- Uses response–returned `usage.prompt_tokens` as the input–token control while allowing small cross–platform accounting variance up to 50 tokens.
- Extends prompt–caching evaluation to cost, throughput, E2E latency, TTFT, provider distribution, reasoning telemetry, and paired statistical tests.
- Scopes every claim to observable response telemetry instead of private routing internals.


2. Experimental Design, Dataset, and Controls

2.1 Dataset Construction

The dataset is `controlled_cache_probe`, covering 4 routing sorts, 2 platforms, 4 groups, and 50 rounds per group. Each round sends identical first/second Chat Completions requests; the second request observes cache-read tokens, TTFT, and E2E latency.

The built-in business templates represent stable long-context workloads such as RAG support, agent tool instructions, marketing automation, and code review.

2.2 Controlled Variables



Controlled-variable rule: within the same `sort/group/round`, both platforms must have first/second `usage.prompt_tokens` deltas no greater than 50 tokens. Total Input Tokens are computed from response-returned usage, not local tokenizer estimates.

2.3 Metric Definitions

Metric	Definition	Direction
Call cache hit rate	Share of second requests with <code>cache_read_tokens > 0</code>	Higher is better
Token cache hit rate	Second-request cache-read tokens / second-request prompt tokens	Higher is better
Observed cost	Sum of first + second request usage/cost values	Lower is better
Throughput	Completion tokens / request latency seconds	Higher is better
E2E latency	Full response latency per request	Lower is better
TTFT	Streaming first-token arrival time	Lower is better
Reasoning telemetry	Reasoning tokens from response usage	Used to explain latency, throughput, and cost

3. Experimental Environment and Data Quality

Item	Configuration
Model	moonshotai/kimi-k2.6
Provider model IDs	infron: moonshotai/kimi-k2.6 ; openrouter: moonshotai/kimi-k2.6
Platforms	Infron and OpenRouter
API protocol	/v1/chat/completions
Routing modes	Throughput First, Price First, Latency First, TTFT First
Groups	4
Rounds per group	50
Workers	24
Request mode	Streaming Chat Completions with TTFT collection
Reasoning / thinking control	No explicit reasoning/thinking parameter; model and platform defaults are preserved
Prompt length tiers	short ≈1500, medium ≈8000, long ≈32000
Excluded records	2

4. Results: Overall Metrics and Main Findings

Routing mode	Platform	Strict pairs	Total Input Tokens	Token cache hit rate	Observed cost	Throughput	E2E latency	Streaming TTFT
Throughput First	Infron	199	6241982	94.91%	\$1.86053300	10.10 tok/s	5479.22 ms	4279.99 ms
Throughput First	OpenRouter	199	6242244	93.76%	\$1.39976252	4.40 tok/s	3599.82 ms	3249.45 ms
Price First	Infron	200	6260562	96.88%	\$1.02713600	2.54 tok/s	5760.20 ms	5137.98 ms
Price First	OpenRouter	200	6260516	87.27%	\$1.39906729	4.74 tok/s	3349.99 ms	2951.81 ms
Latency First	Infron	200	6260437	99.58%	\$1.01120400	3.39 tok/s	4719.13 ms	3894.17 ms
Latency First	OpenRouter	200	6260516	88.05%	\$1.54876474	4.39 tok/s	3609.94 ms	3043.02 ms
TTFT First	Infron	200	6260516	99.90%	\$0.96550800	3.62 tok/s	4425.15 ms	3576.76 ms
TTFT First	OpenRouter	200	6260516	87.97%	\$1.61451600	5.47 tok/s	2898.35 ms	2560.99 ms

4.1 Tail Latency and Statistical Tests

Tail percentiles expose risk that averages hide.

Routing mode	Platform	P50 latency	P95 latency	P99 latency	P50 TTFT	P95 TTFT	P99 TTFT
Throughput First	Infron	4445.00 ms	11184.02 ms	15827.94 ms	3605.06 ms	8473.94 ms	14929.40 ms
Throughput First	OpenRouter	2479.54 ms	7769.27 ms	30620.94 ms	2168.72 ms	6972.79 ms	30493.10 ms
Price First	Infron	3886.19 ms	16688.31 ms	35579.91 ms	3115.13 ms	15290.95 ms	35336.90 ms
Price First	OpenRouter	2838.07 ms	7019.75 ms	12156.43 ms	2416.39 ms	6101.97 ms	10949.25 ms
Latency First	Infron	4003.10 ms	8972.70 ms	11097.75 ms	3257.40 ms	7912.93 ms	9523.80 ms
Latency First	OpenRouter	3048.14 ms	7814.77 ms	13822.59 ms	2634.22 ms	6637.91 ms	8844.16 ms
TTFT First	Infron	3895.71 ms	8039.59 ms	9781.49 ms	3061.59 ms	7098.01 ms	8293.01 ms
TTFT First	OpenRouter	2518.48 ms	5391.63 ms	8474.63 ms	2272.91 ms	4778.55 ms	7189.35 ms

Mean deltas use bootstrap 95% CIs; p-values use paired sign-flip permutation tests.

Routing mode	Metric	Mean delta	95% CI	p-value	Pairs	Interpretation
Throughput First	Latency: OpenRouter – Infron	–3758.79 ms	–4717.17 ms to –2709.97 ms	<0.001	199	Positive means lower Infron latency
Throughput First	TTFT: OpenRouter – Infron	–2061.06 ms	–2931.78 ms to –1054.55 ms	<0.001	199	Positive means lower Infron TTFT
Throughput First	Throughput: Infron – OpenRouter	2.9133 tok/s	1.8454 tok/s to 4.0824 tok/s	<0.001	199	Positive means higher Infron throughput
Throughput First	Cost: OpenRouter – Infron	\$–0.00231543	\$–0.00367426 to \$–0.00101253	<0.001	199	Positive means lower Infron cost
Throughput First	Token Cache Hit: Infron – OpenRouter	–2.81 pp	–6.84 pp to 1.33 pp	0.1637	199	Positive means higher Infron cache hit
Price First	Latency: OpenRouter – Infron	–4820.42 ms	–6383.17 ms to –3480.35 ms	<0.001	200	Positive means lower Infron latency
Price First	TTFT: OpenRouter – Infron	–4372.35 ms	–5927.41 ms to –3032.91 ms	<0.001	200	Positive means lower Infron TTFT
Price First	Throughput: Infron – OpenRouter	–1.8316 tok/s	–2.1415 tok/s to –1.5383 tok/s	<0.001	200	Positive means higher Infron throughput
Price First	Cost: OpenRouter – Infron	\$0.00185966	\$0.00108303 to \$0.00272617	<0.001	200	Positive means lower Infron cost
Price First	Token Cache Hit: Infron – OpenRouter	7.68 pp	3.16 pp to 12.66 pp	<0.001	200	Positive means higher Infron cache hit
Latency First	Latency: OpenRouter – Infron	–2218.37 ms	–2714.07 ms to –1687.13 ms	<0.001	200	Positive means lower Infron latency
Latency First	TTFT: OpenRouter – Infron	–1702.30 ms	–2114.70 ms to –1312.17 ms	<0.001	200	Positive means lower Infron TTFT

Routing mode	Metric	Mean delta	95% CI	p-value	Pairs	Interpretation
Latency First	Throughput: Infron – OpenRouter	-1.4602 tok/s	-1.7083 tok/s to -1.2156 tok/s	<0.001	200	Positive means higher Infron throughput
Latency First	Cost: OpenRouter – Infron	\$0.00268780	\$0.00172956 to \$0.00373797	<0.001	200	Positive means lower Infron cost
Latency First	Token Cache Hit: Infron – OpenRouter	9.21 pp	5.45 pp to 13.35 pp	<0.001	200	Positive means higher Infron cache hit
TTFT First	Latency: OpenRouter – Infron	-3053.60 ms	-3533.09 ms to -2571.02 ms	<0.001	200	Positive means lower Infron latency
TTFT First	TTFT: OpenRouter – Infron	-2031.53 ms	-2416.82 ms to -1663.13 ms	<0.001	200	Positive means lower Infron TTFT
TTFT First	Throughput: Infron – OpenRouter	-2.2692 tok/s	-2.5054 tok/s to -2.0407 tok/s	<0.001	200	Positive means higher Infron throughput
TTFT First	Cost: OpenRouter – Infron	\$0.00324504	\$0.00218847 to \$0.00443978	<0.001	200	Positive means lower Infron cost
TTFT First	Token Cache Hit: Infron – OpenRouter	11.71 pp	7.71 pp to 16.27 pp	<0.001	200	Positive means higher Infron cache hit

4.2 Reasoning / Thinking Control Check

This run does not explicitly set reasoning/thinking parameters and keeps model/platform defaults; this table records reasoning telemetry under default behavior.

Routing mode	Platform	Reasoning tokens	Avg reasoning tokens/request	Reasoning requests	Avg first reasoning token	Avg TTFT	Avg E2E latency
Throughput First	Infron	20836	52.3518	375	4280.08 ms	4279.99 ms	5479.22 ms
Throughput First	OpenRouter	6631	16.6608	398	3249.45 ms	3249.45 ms	3599.82 ms
Price First	Infron	0	0.0000	0	3814.59 ms	5137.98 ms	5760.20 ms
Price First	OpenRouter	6847	17.1175	400	2951.81 ms	2951.81 ms	3349.99 ms
Latency First	Infron	1185	2.9625	79	3980.34 ms	3894.17 ms	4719.13 ms
Latency First	OpenRouter	6903	17.2575	400	3043.02 ms	3043.02 ms	3609.94 ms
TTFT First	Infron	0	0.0000	0	3676.70 ms	3576.76 ms	4425.15 ms
TTFT First	OpenRouter	6997	17.4925	400	2560.99 ms	2560.99 ms	2898.35 ms

4.3 API Protocol Compatibility Matrix

This run uses `/v1/chat/completions` ; this table records success response, usage, cost, and cache telemetry coverage for both platforms under that protocol.

API protocol	Endpoint	Platform	Pairs	Requests	Success rate	Usage coverage	Token usage coverage	Cost coverage	Cache telemetry coverage
chat_completions	/v1/chat/completions	Infron	800	1600	100.00%	100.00%	100.00%	100.00%	100.00%
chat_completions	/v1/chat/completions	OpenRouter	800	1600	100.00%	100.00%	100.00%	100.00%	100.00%

5. Result Visualizations by Routing Mode

The following charts are driven by the strict-paired summary data and present routing-mode, platform, cost, cache, E2E latency, Streaming TTFT, and upstream provider differences.

5.1 Core Metric Chart Overview

The charts below are generated from the same strict-paired summary data.

**Figure 3-E2:
Normalized
capability
contour**

Five metrics are normalized to 0-100, where higher is better.

**Figure 3-E3:
Cost-cache
efficiency plane**

X-axis is token cache hit rate; Y-axis is observed total cost.

**Figure
3-
E1:
Routing-
mode
core
metric
comparison**

Infron and OpenRouter are shown side by side by routing mode.

Figure 3–E4: E2E latency and Streaming TTFT

Solid lines show E2E latency; dashed lines show Streaming TTFT.

Figure 3–E5: Upstream provider distribution

Provider shares explain cache–domain and performance differences.

Routing–Mode Drill–Down Charts

The horizontal axis shows relative advantage: blue to the right indicates Infron advantage, orange to the left indicates OpenRouter advantage.

Figure 4–E1: Throughput First core metric comparison

Direction and magnitude of relative advantage across five metrics.

Figure 4–E2: Price First core metric comparison

Direction and magnitude of relative advantage across five metrics.


Figure 4–E3: Latency First core metric comparison

Direction and magnitude of relative advantage across five metrics.

Figure 4–E4: TTFT First core metric comparison

Direction and magnitude of relative advantage across five metrics.

6. Infron Technical Architecture and Cache/Cost Mechanism



6.1 Multi-Provider Routing and Observable Control Plane

Requests enter a unified API, and the routing layer selects healthy upstream paths according to throughput, price, latency, or ttft objectives. Whether stable long prefixes stay in the same cache domain directly affects second-request cache-read tokens.

6.2 Provider Stick and Cache Affinity

Provider stick is cache affinity, not a permanent provider lock. Its goal is to reduce cache-domain fragmentation while respecting health and SLA constraints.

6.3 Cost-Control Path

Mechanism	Cache-rate effect	Cost effect	Observable signal
Stable prefix detection	Repeated prefixes are more likely to hit cache	Reduces repeated prefill cost	Payload SHA-256 and second-request cache-read tokens
Provider stick / cache affinity	Reduces cross-domain cache fragmentation	Avoids repeated cache warm-up	Provider distribution and token cache hit rate
Health checks and fallback	Protects availability while sometimes sacrificing cache	Reduces failure cost	HTTP status, provider distribution, tail latency
Cost-aware routing	Prefers lower-cost paths under constraints	Reduces total and per-round cost	Observed cost, cost breakdown coverage, cache-read tokens

7. Provider/Route Drill-Down

Routing mode	Platform	Total requests	Attributed requests	Provider distribution
Throughput First	Infron	398	398	novita 306, moonshotai 92

Routing mode	Platform	Total requests	Attributed requests	Provider distribution
Throughput First	OpenRouter	398	398	Together 244, Cloudflare 114, ModelRun 27, Nebius 10, Decart 3
Price First	Infron	400	400	deepinfra 354, alibaba/cn 46
Price First	OpenRouter	400	400	Together 194, Cloudflare 90, Inceptron 90, ModelRun 18, Decart 8
Latency First	Infron	400	400	deepinfra 321, moonshotai 79
Latency First	OpenRouter	400	400	Together 164, Inceptron 71, WandB 70, Cloudflare 59, ModelRun 28, Decart 8
TTFT First	Infron	400	400	deepinfra 400
TTFT First	OpenRouter	400	400	Together 139, WandB 139, Cloudflare 45, Inceptron 43, ModelRun 27, Decart 7

Upstream Provider Detail Distribution

Routing mode	Platform	Upstream provider	Requests	Share	first/second	Covered rounds	Avg TTFT	Avg latency	Prompt tokens	Completion tokens
Throughput First	Infron	novita	306	76.88%	147/159	159	4090.85 ms	5553.31 ms	4944308	20559
Throughput First	Infron	moonshotai	92	23.12%	52/40	52	4909.06 ms	5232.77 ms	1297674	1472
Throughput First	OpenRouter	Together	244	61.31%	126/118	136	1977.85 ms	2225.43 ms	1560099	3904
Throughput First	OpenRouter	Cloudflare	114	28.64%	55/59	68	3494.13 ms	3891.25 ms	3655296	1824
Throughput First	OpenRouter	ModelRun	27	6.78%	10/17	20	13629.44 ms	14771.74 ms	900120	369
Throughput First	OpenRouter	Nebius	10	2.51%	6/4	10	1911.04 ms	2107.15 ms	17640	160
Throughput First	OpenRouter	Decart	3	0.75%	2/1	2	8417.25 ms	8737.09 ms	109089	48
Price First	Infron	deepinfra	354	88.50%	170/184	184	3732.47 ms	4400.81 ms	5604261	5664
Price First	Infron	alibaba/cn	46	11.50%	30/16	30	15954.27 ms	16221.58 ms	656301	184
Price First	OpenRouter	Together	194	48.50%	102/92	106	2109.40 ms	2338.19 ms	973947	3104
Price First	OpenRouter	Cloudflare	90	22.50%	49/41	49	3389.22 ms	3795.48 ms	2864265	1440

Routing mode	Platform	Upstream provider	Requests	Share	first/second	Covered rounds	Avg TTFT	Avg latency	Prompt tokens	Completion tokens
Price First	OpenRouter	Inceptron	90	22.50%	35/55	60	3758.42 ms	4482.13 ms	1593146	1440
Price First	OpenRouter	ModelRun	18	4.50%	12/6	12	4466.40 ms	5030.55 ms	600080	245
Price First	OpenRouter	Decart	8	2.00%	2/6	7	5976.99 ms	6356.27 ms	229078	128
Latency First	Infron	deepinfra	321	80.25%	154/167	167	3751.59 ms	4668.54 ms	5099112	5136
Latency First	Infron	moonshotai	79	19.75%	46/33	46	4473.55 ms	4924.68 ms	1161325	1264
Latency First	OpenRouter	Together	164	41.00%	83/81	85	2146.05 ms	2465.14 ms	753732	2624
Latency First	OpenRouter	Inceptron	71	17.75%	38/33	38	3938.26 ms	5607.09 ms	1211379	1136
Latency First	OpenRouter	WandB	70	17.50%	27/43	47	3904.32 ms	4092.54 ms	1523856	1120
Latency First	OpenRouter	Cloudflare	59	14.75%	34/25	35	3317.52 ms	3788.77 ms	1927601	944
Latency First	OpenRouter	ModelRun	28	7.00%	14/14	15	3148.29 ms	3518.01 ms	553044	388
Latency First	OpenRouter	Decart	8	2.00%	4/4	4	3556.64 ms	4133.92 ms	290904	128
TTFT First	Infron	deepinfra	400	100.00%	200/200	200	3576.76 ms	4425.15 ms	6260516	6400
TTFT First	OpenRouter	Together	139	34.75%	71/68	75	1959.65 ms	2209.50 ms	655767	2224
TTFT First	OpenRouter	WandB	139	34.75%	65/74	82	2564.07 ms	2708.08 ms	2667620	2224
TTFT First	OpenRouter	Cloudflare	45	11.25%	23/22	23	3084.87 ms	3538.23 ms	1500200	720
TTFT First	OpenRouter	Inceptron	43	10.75%	23/20	25	4013.46 ms	5290.99 ms	665707	688
TTFT First	OpenRouter	ModelRun	27	6.75%	14/13	16	2092.10 ms	2227.07 ms	516681	377
TTFT First	OpenRouter	Decart	7	1.75%	4/3	4	3959.51 ms	4133.01 ms	254541	112

7.1 Cache–Rate and Cost Divergence Drill–Down

This table combines cache, cost, provider distribution, and reasoning telemetry to explain routing–mode differences.

Routing mode	Cache–hit delta	Infron cost multiple	Infron top path	OpenRouter top path	Reasoning token delta	Main attribution
Throughput First	+1.15 pp	1.33x	novita 76.88%	Together 61.31%	+14205	Infron has higher cache but higher cost; inspect upstream unit price and completion/ reasoning tokens
Price First	+9.61 pp	0.73x	deepinfra 88.50%	Together 48.50%	–6847	Infron leads both cache and cost
Latency First	+11.52 pp	0.65x	deepinfra 80.25%	Together 41.00%	–5718	Infron leads both cache and cost
TTFT First	+11.93 pp	0.60x	deepinfra 100.00%	Together 34.75%	–6997	Infron leads both cache and cost

8. Stratified Results: Prompt–Length Cache Performance

This section aggregates second–request cache–read tokens, token–level cache hit rate, observed cost, E2E latency, and Streaming TTFT by prompt–length tier. Bold cells mark the advantaged side within each tier.

Prompt–Length Tier Overview

Prompt length tier	Target tokens	Platform	Pairs	Second prompt tokens	Second cache read tokens	Token cache hit rate	Observed cost	Avg E2E latency	Avg TTFT
short	1500	Infron	268	472703	458445	96.98%	\$0.22535400	3543.85 ms	2691.20 ms
short	1500	OpenRouter	268	472752	431552	91.29%	\$0.28809660	2213.64 ms	1861.27 ms
medium	8000	Infron	267	2439265	2338256	95.86%	\$1.01142000	4450.15 ms	3591.57 ms
medium	8000	OpenRouter	267	2439312	2245344	92.05%	\$1.22135368	2997.84 ms	2634.75 ms
long	32000	Infron	264	9599782	9442026	98.36%	\$3.62760700	7323.16 ms	6414.05 ms
long	32000	OpenRouter	264	9599832	8491200	88.45%	\$4.45266027	4902.81 ms	4376.93 ms

Prompt Length x Routing Mode Cache Hit Rate

Prompt length tier	Routing mode	Infron	OpenRouter	Winner
short	Throughput First	92.31%	96.50%	OpenRouter
short	Price First	97.01%	90.74%	Infron
short	Latency First	99.41%	92.04%	Infron
short	TTFT First	99.20%	85.86%	Infron
medium	Throughput First	88.54%	96.46%	OpenRouter
medium	Price First	96.81%	90.64%	Infron
medium	Latency First	98.16%	90.59%	Infron
medium	TTFT First	99.82%	90.57%	Infron
long	Throughput First	96.64%	92.95%	Infron
long	Price First	96.89%	86.25%	Infron
long	Latency First	99.95%	87.21%	Infron
long	TTFT First	99.95%	87.41%	Infron

9. Stratified Results: Group–Level Stability

Throughput First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	50	50	91.49%	\$0.70589100	15313.44 ms	14903.04 ms
Infron	2	50	50	92.36%	\$0.46157600	10518.61 ms	8090.84 ms
Infron	3	49	49	98.49%	\$0.37659400	10438.36 ms	8074.49 ms
Infron	4	50	50	97.33%	\$0.31647200	9682.92 ms	6747.94 ms
OpenRouter	1	50	50	85.27%	\$0.42537508	5543.31 ms	5505.53 ms
OpenRouter	2	50	50	94.93%	\$0.36337632	7954.02 ms	7927.13 ms
OpenRouter	3	49	49	99.65%	\$0.31050348	10795.44 ms	8942.72 ms
OpenRouter	4	50	50	94.99%	\$0.30050764	4798.96 ms	4653.01 ms

Price First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	50	50	92.42%	\$0.25295000	30554.29 ms	30250.22 ms
Infron	2	50	50	99.87%	\$0.26949900	7895.14 ms	6792.94 ms
Infron	3	50	50	95.29%	\$0.26734900	8592.21 ms	7717.53 ms

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	4	50	50	99.86%	\$0.23733800	8877.56 ms	7865.57 ms
OpenRouter	1	50	50	86.31%	\$0.37109449	6897.29 ms	6449.71 ms
OpenRouter	2	50	50	92.78%	\$0.34899592	5808.95 ms	5020.80 ms
OpenRouter	3	50	50	76.27%	\$0.39277182	9044.22 ms	7899.05 ms
OpenRouter	4	50	50	93.88%	\$0.28620506	5549.87 ms	5199.22 ms

Latency First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	50	50	98.60%	\$0.27559400	9149.24 ms	8846.74 ms
Infron	2	50	50	99.90%	\$0.25370300	9149.63 ms	7729.14 ms
Infron	3	50	50	99.89%	\$0.24500600	9173.18 ms	7089.74 ms
Infron	4	50	50	99.89%	\$0.23690100	8222.10 ms	6998.55 ms
OpenRouter	1	50	50	78.45%	\$0.47371294	7337.08 ms	6799.77 ms
OpenRouter	2	50	50	88.89%	\$0.40170657	8410.81 ms	7676.65 ms
OpenRouter	3	50	50	86.34%	\$0.34924236	8048.15 ms	6064.69 ms
OpenRouter	4	50	50	98.55%	\$0.32410287	6717.82 ms	5769.41 ms

TTFT First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	50	50	99.87%	\$0.23711000	7834.03 ms	6771.93 ms
Infron	2	50	50	99.89%	\$0.24706100	7927.66 ms	7021.17 ms
Infron	3	50	50	99.91%	\$0.24475900	8468.28 ms	7389.17 ms
Infron	4	50	50	99.91%	\$0.23657800	7806.07 ms	6875.62 ms
OpenRouter	1	50	50	81.50%	\$0.45947412	5223.55 ms	4695.71 ms
OpenRouter	2	50	50	92.24%	\$0.43593146	5956.50 ms	5236.29 ms
OpenRouter	3	50	50	84.49%	\$0.37611854	5772.09 ms	5250.33 ms
OpenRouter	4	50	50	93.58%	\$0.34299188	4317.41 ms	3860.15 ms

10. Discussion: Business Value, Boundaries, and Engineering Implications

Business decisions should not rely on one metric. Stable long-context and high-frequency template workloads should prioritize cache rate and cost; realtime interaction must constrain TTFT and E2E latency; batch processing often prioritizes throughput and failure cost.

Routing mode	Business objective	Observed result	Scenarios	Caveat
Throughput First	Maximize output capacity per unit time	OpenRouter leads overall (3/5 metrics)	Batch generation, offline summaries, backend processing	First-token and E2E response are faster; check cache and cost
Price First	Minimize request and token cost	OpenRouter leads overall (3/5 metrics)	High-frequency templates, support automation, marketing, RAG prefixes	Cache and cost are stronger; still check speed SLA
Latency First	Minimize full-response waiting time	OpenRouter leads overall (3/5 metrics)	Online chat, agent chains, IDE/writing assistants, realtime tools	Cache and cost are stronger; still check speed SLA
TTFT First	Minimize streaming first-token time	OpenRouter leads overall (3/5 metrics)	Streaming chat, realtime copilots, first-screen feedback	Cache and cost are stronger; still check speed SLA

11. Conclusion

Routing-Mode Conclusions

Routing mode	Objective	Cache winner	Cost winner	Throughput winner	E2E latency winner	TTFT winner	Interpretation
Throughput First	Maximize output capacity per unit time	Infron	OpenRouter	Infron	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)
Price First	Minimize request and token cost	Infron	Infron	OpenRouter	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)
Latency First	Minimize full-response waiting time	Infron	Infron	OpenRouter	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)
TTFT First	Minimize streaming first-token time	Infron	Infron	OpenRouter	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)

12. Limitations, Missing Data, and Next Experiments

Limitation	Impact	Next step	Current handling
Full routing trace	Cannot prove every provider choice, fallback, and retry path hop by hop	Add provider routing trace, decision logs, and fallback reasons	Use only returned provider fields and provider distribution
Longer time window	4x50 observes short-window stability but not day-level drift	Add soak tests and repeated windows	Scope conclusions to this run
Production corpus	Built-in templates do not cover every workload distribution	Use sanitized production-stratified corpora	Discuss representative long-context templates only

Limitation	Impact	Next step	Current handling
Cost-field consistency	Cost coverage and semantics may differ by platform	Reconcile with billing and provider cost breakdown	Use only explicitly returned cost fields

13. Reproducibility Appendix

Artifact	Path
Summary	export/kimi_k26_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions summary.json
Paired dataset	export/kimi_k26_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions benchmark_pairs.csv
Request- level dataset	export/kimi_k26_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions benchmark_requests.jsonl
Filtered structured records	export/kimi_k26_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions records.json
Excluded- record audit	export/kimi_k26_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions records_excluded.json
Test source	tests/test_rerun_routing_sort_cache_cost_ab.py
Benchmark runner source	scripts/rerun_routing_sort_cache_cost_ab.py
HTML report renderer source	scripts/render_glm52_deepseek_style_report.py
Dataset reference	business_representative built-in representative business templates; request-level export is benchmark_requests