

Infron vs OpenRouter Routing, Cache, Cost, and Streaming Performance A/B Test Report

Abstract and Executive Outline

Keywords: Prompt Caching; A/B Testing; Provider Routing; Cache Affinity; Latency; Throughput; Cost Attribution; qwen3.7-flash

Abstract

This report evaluates `qwen/qwen3.7-flash` on Infron and OpenRouter across cache reuse, observed cost, throughput, E2E latency, and Streaming TTFT under Prompt Caching workloads.

The main findings are: Infron leads Cache hit rate in 1/4 routing modes; OpenRouter leads in 3/4; 0/4 tie; Infron leads Observed cost in all routing modes; Infron leads Throughput in all routing modes; OpenRouter leads E2E latency in all routing modes; Infron leads Streaming TTFT in 1/4 routing modes; OpenRouter leads in 3/4; 0/4 tie.

Overall, Infron's cross-mode strengths are throughput, observed cost, while OpenRouter's cross-mode strengths are Streaming TTFT, E2E latency, cache reuse. Platform choice should be driven by workload objective rather than a single headline metric.

Figure 0: Normalized Capability Radar

The five radar axes represent throughput, price, E2E latency, Streaming TTFT, and cache hit rate. Every metric is normalized to a 0–100 score, with farther outward meaning better.

Thick solid lines show platform–level contours; translucent lines and points show individual routing modes.

Conclusion Overview: Core Metric Winners by Routing Mode

Based on strict A/B paired samples. Blue represents Infron, orange represents OpenRouter; gold cells mark the winner for the routing objective.

Throughput Infron wins 4/4 Max advantage 12.24%, higher is better	Observed cost Infron wins 4/4 Max advantage 75.91%, lower is better	E2E latency OpenRouter wins 4/4 Max advantage 97.35%, lower is better	Streaming TTFT OpenRouter wins 3/4 Max advantage 4.02%, lower is better	Cache hit rate OpenRouter wins 3/4 Max advantage 25.07%, higher is better
--	--	--	--	--

Routing mode	Throughput objective	Cost objective	Latency objective	TTFT objective	Cache result
Throughput First throughput	Infron advantage 10.97%	Infron advantage 37.92%	OpenRouter advantage 95.81%	OpenRouter advantage 4.02%	OpenRouter advantage 25.07%
Price First price	Infron advantage 12.24%	Infron advantage 67.92%	OpenRouter advantage 82.73%	Infron advantage 1.19%	OpenRouter advantage 3.92%
Latency First latency	Infron advantage 11.57%	Infron advantage 68.27%	OpenRouter advantage 97.35%	OpenRouter advantage 3.40%	Infron advantage 0.23%
TTFT First ttft	Infron advantage 10.93%	Infron advantage 75.91%	OpenRouter advantage 96.15%	OpenRouter advantage 1.43%	OpenRouter advantage 4.88%

Executive Outline

Dimension	Conclusion	Evidence
Controls	First/second <code>usage.prompt_tokens</code> deltas are limited to 50 tokens within each <code>sort/group/round</code> pair.	Methods and data quality
Cache reuse	Infron leads Cache hit rate in 1/4 routing modes; OpenRouter leads in 3/4; 0/4 tie	Overall metrics and mechanism section
Observed cost	Infron leads Observed cost in all routing modes	Overall metrics and provider drill-down
Performance	Infron leads Throughput in all routing modes; OpenRouter leads E2E latency in all routing modes; Infron leads Streaming TTFT in 1/4 routing modes; OpenRouter leads in 3/4; 0/4 tie	Charts and statistical tests
Attribution boundary	Claims use observable response telemetry: usage, cost, TTFT, latency, provider fields, and cache tokens.	Provider/Route drill-down
Business meaning	Long-context, RAG-prefix, agent-tool, and high-frequency template workloads should evaluate cache rate, cost, first-token latency, and E2E latency together.	Discussion and conclusion

Routing–Mode Conclusions

Routing mode	Objective	Cache winner	Cost winner	Throughput winner	E2E latency winner	TTFT winner	Interpretation
Throughput First	Maximize output capacity per unit time	OpenRouter	Infron	Infron	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)
Price First	Minimize request and token cost	OpenRouter	Infron	Infron	OpenRouter	Infron	Infron leads overall (3/5 metrics)
Latency First	Minimize full–response waiting time	Infron	Infron	Infron	OpenRouter	OpenRouter	Infron leads overall (3/5 metrics)
TTFT First	Minimize streaming first–token time	OpenRouter	Infron	Infron	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)

1. Introduction: Background, Questions, and Contributions

LLM inference–platform behavior is shaped by provider routing, prompt caching, streaming response handling, cost attribution, and fallback policy. This report treats the platform as an observable system and measures speed, cost, cache reuse, and first–token experience through strict A/B pairs.

1.1 Research Hypotheses

Hypothesis	Statement	Validation metric
H1	Stronger provider/cache affinity improves token–level cache hit rate for repeated stable prefixes.	Second–request cache–read tokens and token cache hit rate
H2	Higher cache hit rate can reduce observed cost, but does not necessarily reduce TTFT or E2E latency.	Observed cost, average TTFT, average request latency
H3	Different routing sorts change provider selection and produce different cost/throughput/latency frontiers.	Provider distribution, throughput, latency, cost

1.2 Contributions

- Uses response–returned `usage.prompt_tokens` as the input–token control while allowing small cross–platform accounting variance up to 50 tokens.
- Extends prompt–caching evaluation to cost, throughput, E2E latency, TTFT, provider distribution, reasoning telemetry, and paired statistical tests.
- Scopes every claim to observable response telemetry instead of private routing internals.

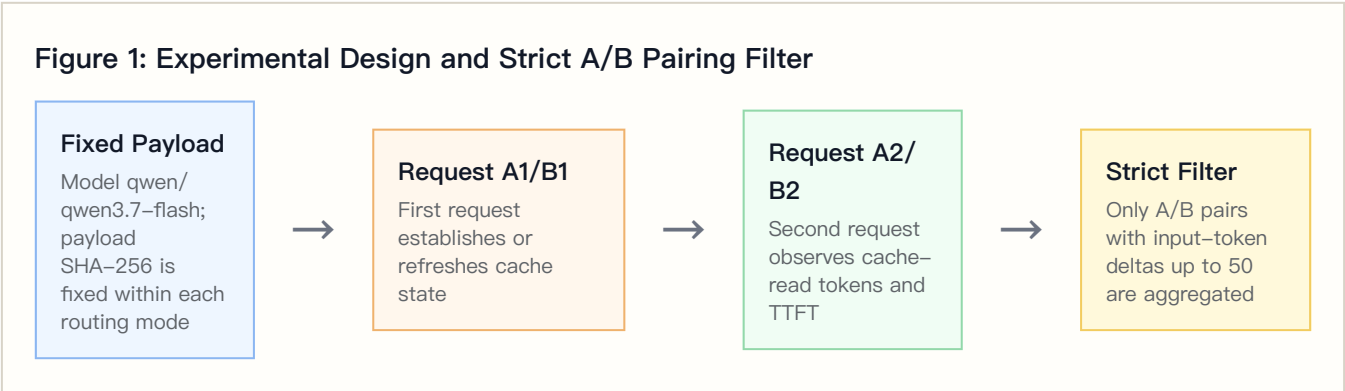
2. Experimental Design, Dataset, and Controls

2.1 Dataset Construction

The dataset is `controlled_cache_probe`, covering 4 routing sorts, 2 platforms, 4 groups, and 50 rounds per group. Each round sends identical first/second Chat Completions requests; the second request observes cache-read tokens, TTFT, and E2E latency.

The built-in business templates represent stable long-context workloads such as RAG support, agent tool instructions, marketing automation, and code review.

2.2 Controlled Variables



Controlled-variable rule: within the same `sort/group/round`, both platforms must have first/second `usage.prompt_tokens` deltas no greater than 50 tokens. Total Input Tokens are computed from response-returned usage, not local tokenizer estimates.

2.3 Metric Definitions

Metric	Definition	Direction
Call cache hit rate	Share of second requests with <code>cache_read_tokens > 0</code>	Higher is better
Token cache hit rate	Second-request cache-read tokens / second-request prompt tokens	Higher is better
Observed cost	Sum of first + second request usage/cost values	Lower is better
Throughput	Completion tokens / request latency seconds	Higher is better
E2E latency	Full response latency per request	Lower is better
TTFT	Streaming first-token arrival time	Lower is better
Reasoning telemetry	Reasoning tokens from response usage	Used to explain latency, throughput, and cost

3. Experimental Environment and Data Quality

Item	Configuration
Model	qwen/qwen3.7-flash
Provider model IDs	infron: qwen/qwen3.7-flash ; openrouter: qwen/qwen3.7-flash
Platforms	Infron and OpenRouter
API protocol	/v1/chat/completions
Routing modes	Throughput First, Price First, Latency First, TTFT First
Groups	4
Rounds per group	50
Workers	24
Request mode	Streaming Chat Completions with TTFT collection
Reasoning / thinking control	No explicit reasoning/thinking parameter; model and platform defaults are preserved
Prompt length tiers	short ≈1500, medium ≈8000, long ≈32000
Excluded records	230

4. Results: Overall Metrics and Main Findings

Routing mode	Platform	Strict pairs	Total Input Tokens	Token cache hit rate	Observed cost	Throughput	E2E latency	Streaming TTFT
Throughput First	Infron	163	5667484	53.89%	\$0.14534700	71.44 tok/s	5906.04 ms	3006.71 ms
Throughput First	OpenRouter	163	5667484	67.40%	\$0.20046654	5.97 tok/s	3016.26 ms	2890.54 ms
Price First	Infron	178	6497986	67.46%	\$0.12428800	69.26 tok/s	6287.51 ms	3247.79 ms
Price First	OpenRouter	178	6497986	70.11%	\$0.20870199	5.23 tok/s	3440.89 ms	3286.44 ms
Latency First	Infron	177	6086188	69.29%	\$0.11321900	73.55 tok/s	6072.71 ms	3035.85 ms
Latency First	OpenRouter	177	6086188	69.13%	\$0.19051714	5.85 tok/s	3077.13 ms	2935.90 ms
TTFT First	Infron	167	5711636	59.02%	\$0.11308200	68.65 tok/s	6134.60 ms	3031.17 ms
TTFT First	OpenRouter	167	5711636	61.90%	\$0.19892656	5.75 tok/s	3127.46 ms	2988.47 ms

4.1 Tail Latency and Statistical Tests

Tail percentiles expose risk that averages hide.

Routing mode	Platform	P50 latency	P95 latency	P99 latency	P50 TTFT	P95 TTFT	P99 TTFT
Throughput First	Infron	5699.36 ms	8905.14 ms	10907.52 ms	2566.90 ms	5920.58 ms	7224.64 ms
Throughput First	OpenRouter	2837.99 ms	5214.80 ms	6625.83 ms	2720.70 ms	5109.37 ms	6027.52 ms
Price First	Infron	6031.07 ms	9032.75 ms	12245.83 ms	3066.21 ms	5959.19 ms	7922.48 ms
Price First	OpenRouter	3201.89 ms	6689.85 ms	8921.92 ms	3076.23 ms	6177.10 ms	8866.79 ms
Latency First	Infron	5872.34 ms	8723.91 ms	9531.46 ms	2841.24 ms	5575.12 ms	6260.39 ms
Latency First	OpenRouter	2969.81 ms	5456.74 ms	7410.76 ms	2851.19 ms	5340.58 ms	7206.51 ms
TTFT First	Infron	5929.55 ms	8870.69 ms	10394.01 ms	2977.10 ms	5491.02 ms	6103.65 ms
TTFT First	OpenRouter	3014.78 ms	6096.69 ms	8425.05 ms	2844.58 ms	5551.33 ms	8139.81 ms

Mean deltas use bootstrap 95% CIs; p-values use paired sign-flip permutation tests.

Routing mode	Metric	Mean delta	95% CI	p-value	Pairs	Interpretation
Throughput First	Latency: OpenRouter – Infron	–5779.54 ms	–6086.35 ms to –5471.26 ms	<0.001	163	Positive means lower Infron latency
Throughput First	TTFT: OpenRouter – Infron	–232.34 ms	–453.01 ms to –0.10 ms	0.0500	163	Positive means lower Infron TTFT
Throughput First	Throughput: Infron – OpenRouter	66.0393 tok/s	63.6555 tok/s to 68.2486 tok/s	<0.001	163	Positive means higher Infron throughput
Throughput First	Cost: OpenRouter – Infron	\$0.00033816	\$0.00018829 to \$0.00049834	<0.001	163	Positive means lower Infron cost
Throughput First	Token Cache Hit: Infron – OpenRouter	–5.32 pp	–12.32 pp to 1.80 pp	0.1492	163	Positive means higher Infron cache hit
Price First	Latency: OpenRouter – Infron	–5693.25 ms	–6166.21 ms to –5252.16 ms	<0.001	178	Positive means lower Infron latency
Price First	TTFT: OpenRouter – Infron	77.30 ms	–189.34 ms to 346.93 ms	0.5714	178	Positive means lower Infron TTFT
Price First	Throughput: Infron – OpenRouter	63.7745 tok/s	61.9940 tok/s to 65.4792 tok/s	<0.001	178	Positive means higher Infron throughput
Price First	Cost: OpenRouter – Infron	\$0.00047424	\$0.00031208 to \$0.00064485	<0.001	178	Positive means lower Infron cost
Price First	Token Cache Hit: Infron – OpenRouter	1.70 pp	–4.45 pp to 8.44 pp	0.6383	178	Positive means higher Infron cache hit
Latency First	Latency: OpenRouter – Infron	–5991.16 ms	–6341.65 ms to –5632.86 ms	<0.001	177	Positive means lower Infron latency
Latency First	TTFT: OpenRouter – Infron	–199.90 ms	–412.01 ms to 14.55 ms	0.0650	177	Positive means lower Infron TTFT

Routing mode	Metric	Mean delta	95% CI	p-value	Pairs	Interpretation
Latency First	Throughput: Infron – OpenRouter	67.6055 tok/s	65.6004 tok/s to 69.3930 tok/s	<0.001	177	Positive means higher Infron throughput
Latency First	Cost: OpenRouter – Infron	\$0.00043671	\$0.00030013 to \$0.00058454	<0.001	177	Positive means lower Infron cost
Latency First	Token Cache Hit: Infron – OpenRouter	-1.34 pp	-7.30 pp to 4.31 pp	0.7091	177	Positive means higher Infron cache hit
TTFT First	Latency: OpenRouter – Infron	-6014.30 ms	-6454.41 ms to -5574.72 ms	<0.001	167	Positive means lower Infron latency
TTFT First	TTFT: OpenRouter – Infron	-85.40 ms	-337.75 ms to 171.51 ms	0.5244	167	Positive means lower Infron TTFT
TTFT First	Throughput: Infron – OpenRouter	62.7361 tok/s	61.0764 tok/s to 64.4892 tok/s	<0.001	167	Positive means higher Infron throughput
TTFT First	Cost: OpenRouter – Infron	\$0.00051404	\$0.00033418 to \$0.00070593	<0.001	167	Positive means lower Infron cost
TTFT First	Token Cache Hit: Infron – OpenRouter	-3.07 pp	-9.35 pp to 3.77 pp	0.4054	167	Positive means higher Infron cache hit

4.2 Reasoning / Thinking Control Check

This run does not explicitly set reasoning/thinking parameters and keeps model/platform defaults; this table records reasoning telemetry under default behavior.

Routing mode	Platform	Reasoning tokens	Avg reasoning tokens/request	Reasoning requests	Avg first reasoning token	Avg TTFT	Avg E2E latency
Throughput First	Infron	134940	413.9264	326	3036.10 ms	3006.71 ms	5906.04 ms
Throughput First	OpenRouter	5216	16.0000	326	2890.54 ms	2890.54 ms	3016.26 ms
Price First	Infron	152191	427.5028	356	3274.56 ms	3247.79 ms	6287.51 ms
Price First	OpenRouter	5696	16.0000	356	3286.44 ms	3286.44 ms	3440.89 ms
Latency First	Infron	155291	438.6751	354	3060.04 ms	3035.85 ms	6072.71 ms
Latency First	OpenRouter	5664	16.0000	354	2935.90 ms	2935.90 ms	3077.13 ms
TTFT First	Infron	137991	413.1467	334	3059.04 ms	3031.17 ms	6134.60 ms
TTFT First	OpenRouter	5344	16.0000	334	2988.47 ms	2988.47 ms	3127.46 ms

4.3 API Protocol Compatibility Matrix

This run uses `/v1/chat/completions` ; this table records success response, usage, cost, and cache telemetry coverage for both platforms under that protocol.

API protocol	Endpoint	Platform	Pairs	Requests	Success rate	Usage coverage	Token usage coverage	Cost coverage	Cache telemetry coverage
chat_completions	/v1/chat/completions	Infron	800	1600	100.00%	100.00%	100.00%	100.00%	100.00%
chat_completions	/v1/chat/completions	OpenRouter	800	1600	92.81%	100.00%	100.00%	100.00%	100.00%

5. Result Visualizations by Routing Mode

The following charts are driven by the strict-paired summary data and present routing-mode, platform, cost, cache, E2E latency, Streaming TTFT, and upstream provider differences.

5.1 Core Metric Chart Overview

The charts below are generated from the same strict-paired summary data.

**Figure 3-E2:
Normalized
capability
contour**

Five metrics are normalized to 0-100, where higher is better.

**Figure 3-E3:
Cost-cache
efficiency plane**

X-axis is token cache hit rate; Y-axis is observed total cost.

**Figure 3-E1:
Routing-mode
core
metric
comparison**

Infron and OpenRouter are shown side by side by routing mode.

Figure 3–E4: E2E latency and Streaming TTFT

Solid lines show E2E latency; dashed lines show Streaming TTFT.

Figure 3–E5: Upstream provider distribution

Provider shares explain cache–domain and performance differences.

Routing–Mode Drill–Down Charts

The horizontal axis shows relative advantage: blue to the right indicates Infron advantage, orange to the left indicates OpenRouter advantage.

Figure 4–E1: Throughput First core metric comparison

Direction and magnitude of relative advantage across five metrics.

Figure 4–E2: Price First core metric comparison

Direction and magnitude of relative advantage across five metrics.


Figure 4–E3: Latency First core metric comparison

Direction and magnitude of relative advantage across five metrics.

Figure 4–E4: TTFT First core metric comparison

Direction and magnitude of relative advantage across five metrics.

6. Infron Technical Architecture and Cache/Cost Mechanism



6.1 Multi-Provider Routing and Observable Control Plane

Requests enter a unified API, and the routing layer selects healthy upstream paths according to throughput, price, latency, or ttft objectives. Whether stable long prefixes stay in the same cache domain directly affects second-request cache-read tokens.

6.2 Provider Stick and Cache Affinity

Provider stick is cache affinity, not a permanent provider lock. Its goal is to reduce cache-domain fragmentation while respecting health and SLA constraints.

6.3 Cost-Control Path

Mechanism	Cache-rate effect	Cost effect	Observable signal
Stable prefix detection	Repeated prefixes are more likely to hit cache	Reduces repeated prefill cost	Payload SHA-256 and second-request cache-read tokens
Provider stick / cache affinity	Reduces cross-domain cache fragmentation	Avoids repeated cache warm-up	Provider distribution and token cache hit rate
Health checks and fallback	Protects availability while sometimes sacrificing cache	Reduces failure cost	HTTP status, provider distribution, tail latency
Cost-aware routing	Prefers lower-cost paths under constraints	Reduces total and per-round cost	Observed cost, cost breakdown coverage, cache-read tokens

7. Provider/Route Drill-Down

Routing mode	Platform	Total requests	Attributed requests	Provider distribution
Throughput First	Infron	326	326	alibaba/jp 326
Throughput First	OpenRouter	326	326	Alibaba 326
Price First	Infron	356	356	alibaba/jp 356

Routing mode	Platform	Total requests	Attributed requests	Provider distribution
Price First	OpenRouter	356	356	Alibaba 356
Latency First	Infron	354	354	alibaba/jp 354
Latency First	OpenRouter	354	354	Alibaba 354
TTFT First	Infron	334	334	alibaba/jp 334
TTFT First	OpenRouter	334	334	Alibaba 334

Upstream Provider Detail Distribution

Routing mode	Platform	Upstream provider	Requests	Share	first/second	Covered rounds	Avg TTFT	Avg latency	Prompt tokens	Completion tokens	Reasoning tokens
Throughput First	Infron	alibaba/jp	326	100.00%	163/163	163	3006.71 ms	5906.04 ms	5667484	137546	134
Throughput First	OpenRouter	Alibaba	326	100.00%	163/163	163	2890.54 ms	3016.26 ms	5667484	5868	52
Price First	Infron	alibaba/jp	356	100.00%	178/178	178	3247.79 ms	6287.51 ms	6497986	155037	152
Price First	OpenRouter	Alibaba	356	100.00%	178/178	178	3286.44 ms	3440.89 ms	6497986	6408	56
Latency First	Infron	alibaba/jp	354	100.00%	177/177	177	3035.85 ms	6072.71 ms	6086188	158119	155
Latency First	OpenRouter	Alibaba	354	100.00%	177/177	177	2935.90 ms	3077.13 ms	6086188	6372	56
TTFT First	Infron	alibaba/jp	334	100.00%	167/167	167	3031.17 ms	6134.60 ms	5711636	140661	137
TTFT First	OpenRouter	Alibaba	334	100.00%	167/167	167	2988.47 ms	3127.46 ms	5711636	6012	53

7.1 Cache–Rate and Cost Divergence Drill–Down

This table combines cache, cost, provider distribution, and reasoning telemetry to explain routing–mode differences.

Routing mode	Cache–hit delta	Infron cost multiple	Infron top path	OpenRouter top path	Reasoning token delta	Main attribution
Throughput First	–13.51 pp	0.73x	alibaba/jp 100.00%	Alibaba 100.00%	+129724	Cache and cost move in different directions; evaluate with speed metrics
Price First	–2.65 pp	0.60x	alibaba/jp 100.00%	Alibaba 100.00%	+146495	Cache and cost move in different directions; evaluate with speed metrics

Routing mode	Cache-hit delta	Infron cost multiple	Infron top path	OpenRouter top path	Reasoning token delta	Main attribution
Latency First	+0.16 pp	0.59x	alibaba/jp 100.00%	Alibaba 100.00%	+149627	Infron leads both cache and cost
TTFT First	−2.88 pp	0.57x	alibaba/jp 100.00%	Alibaba 100.00%	+132647	Cache and cost move in different directions; evaluate with speed metrics

8. Stratified Results: Prompt–Length Cache Performance

This section aggregates second-request cache-read tokens, token-level cache hit rate, observed cost, E2E latency, and Streaming TTFT by prompt-length tier. Bold cells mark the advantaged side within each tier.

Prompt–Length Tier Overview

Prompt length tier	Target tokens	Platform	Pairs	Second prompt tokens	Second cache read tokens	Token cache hit rate	Observed cost	Avg E2E latency	Avg TTFT
short	1500	Infron	231	460111	0	0.00%	\$0.02354300	5117.83 ms	1788.12 ms
short	1500	OpenRouter	231	460111	0	0.00%	\$0.02868774	2068.23 ms	1915.67 ms
medium	8000	Infron	232	2395822	697728	29.12%	\$0.06336200	6656.13 ms	2889.94 ms
medium	8000	OpenRouter	232	2395822	695680	29.04%	\$0.11035802	3182.88 ms	3038.52 ms
long	32000	Infron	222	9125714	6815232	74.68%	\$0.40903100	6553.02 ms	4631.66 ms
long	32000	OpenRouter	222	9125714	7363584	80.69%	\$0.65956648	4301.25 ms	4177.56 ms

Prompt Length x Routing Mode Cache Hit Rate

Prompt length tier	Routing mode	Infron	OpenRouter	Winner
short	Throughput First	0.00%	0.00%	tie
short	Price First	0.00%	0.00%	tie
short	Latency First	0.00%	0.00%	tie
short	TTFT First	0.00%	0.00%	tie
medium	Throughput First	46.51%	45.01%	Infron
medium	Price First	30.84%	19.63%	Infron

Prompt length tier	Routing mode	Infron	OpenRouter	Winner
medium	Latency First	24.23%	29.77%	OpenRouter
medium	TTFT First	14.25%	21.99%	OpenRouter
long	Throughput First	58.64%	76.96%	OpenRouter
long	Price First	79.66%	85.91%	OpenRouter
long	Latency First	85.07%	83.37%	Infron
long	TTFT First	73.71%	75.51%	OpenRouter

9. Stratified Results: Group–Level Stability

Throughput First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	48	48	52.96%	\$0.04762900	8666.08 ms	5683.02 ms
Infron	2	30	30	45.17%	\$0.02461800	7759.77 ms	6049.56 ms
Infron	3	44	44	57.45%	\$0.03744000	9870.37 ms	6103.58 ms
Infron	4	41	41	57.93%	\$0.03566000	8238.86 ms	5938.17 ms
OpenRouter	1	48	48	67.37%	\$0.07207349	4993.49 ms	4873.86 ms
OpenRouter	2	30	30	73.60%	\$0.03168526	5560.44 ms	5493.56 ms
OpenRouter	3	44	44	57.96%	\$0.05789421	5840.77 ms	5355.73 ms
OpenRouter	4	41	41	72.59%	\$0.03881359	4670.82 ms	4629.94 ms

Price First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	50	50	69.92%	\$0.03242800	8868.98 ms	5958.35 ms
Infron	2	37	37	82.69%	\$0.02072400	8286.73 ms	5139.67 ms
Infron	3	49	49	53.45%	\$0.03863600	9370.96 ms	6172.76 ms
Infron	4	42	42	67.35%	\$0.03250000	10016.01 ms	6200.28 ms
OpenRouter	1	50	50	67.75%	\$0.05904559	5330.42 ms	5267.33 ms
OpenRouter	2	37	37	63.12%	\$0.04862474	5749.84 ms	5587.03 ms
OpenRouter	3	49	49	71.49%	\$0.05522492	7698.67 ms	7418.68 ms
OpenRouter	4	42	42	77.20%	\$0.04580675	7145.65 ms	6966.98 ms

Latency First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	44	44	71.92%	\$0.02755200	8759.05 ms	5893.91 ms
Infron	2	45	45	67.47%	\$0.03075500	7965.10 ms	5486.28 ms
Infron	3	41	41	60.29%	\$0.02772500	9171.96 ms	5379.96 ms
Infron	4	47	47	76.84%	\$0.02718700	8400.30 ms	4580.27 ms
OpenRouter	1	44	44	56.18%	\$0.04311926	5234.62 ms	5153.55 ms
OpenRouter	2	45	45	71.13%	\$0.05412047	5674.84 ms	5609.34 ms
OpenRouter	3	41	41	77.61%	\$0.04212614	4978.07 ms	4414.54 ms
OpenRouter	4	47	47	70.54%	\$0.05115127	4818.79 ms	4746.44 ms

TTFT First

Platform	Group	Rounds	Successful rounds	Token cache hit rate	Observed cost	P95 latency	P95 TTFT
Infron	1	47	47	57.88%	\$0.03227800	8318.36 ms	5404.08 ms
Infron	2	38	38	47.51%	\$0.02506900	8721.57 ms	5370.84 ms
Infron	3	38	38	66.97%	\$0.02551800	9881.11 ms	5421.53 ms
Infron	4	44	44	62.19%	\$0.03021700	9091.01 ms	5708.10 ms
OpenRouter	1	47	47	62.75%	\$0.05790167	5056.57 ms	4933.56 ms
OpenRouter	2	38	38	77.32%	\$0.03294962	4261.46 ms	4207.43 ms
OpenRouter	3	38	38	54.10%	\$0.05359975	7012.86 ms	6867.62 ms
OpenRouter	4	44	44	55.54%	\$0.05447553	5279.44 ms	5155.22 ms

10. Discussion: Business Value, Boundaries, and Engineering Implications

Business decisions should not rely on one metric. Stable long-context and high-frequency template workloads should prioritize cache rate and cost; realtime interaction must constrain TTFT and E2E latency; batch processing often prioritizes throughput and failure cost.

Routing mode	Business objective	Observed result	Scenarios	Caveat
Throughput First	Maximize output capacity per unit time	OpenRouter leads overall (3/5 metrics)	Batch generation, offline summaries, backend processing	First-token and E2E response are faster; check cache and cost
Price First	Minimize request and token cost	Infron leads overall (3/5 metrics)	High-frequency templates, support automation, marketing, RAG prefixes	Decide with budget, SLA, and cache stability together

Routing mode	Business objective	Observed result	Scenarios	Caveat
Latency First	Minimize full-response waiting time	Infron leads overall (3/5 metrics)	Online chat, agent chains, IDE/writing assistants, realtime tools	Cache and cost are stronger; still check speed SLA
TTFT First	Minimize streaming first-token time	OpenRouter leads overall (3/5 metrics)	Streaming chat, realtime copilots, first-screen feedback	First-token and E2E response are faster; check cache and cost

11. Conclusion

Routing-Mode Conclusions

Routing mode	Objective	Cache winner	Cost winner	Throughput winner	E2E latency winner	TTFT winner	Interpretation
Throughput First	Maximize output capacity per unit time	OpenRouter	Infron	Infron	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)
Price First	Minimize request and token cost	OpenRouter	Infron	Infron	OpenRouter	Infron	Infron leads overall (3/5 metrics)
Latency First	Minimize full-response waiting time	Infron	Infron	Infron	OpenRouter	OpenRouter	Infron leads overall (3/5 metrics)
TTFT First	Minimize streaming first-token time	OpenRouter	Infron	Infron	OpenRouter	OpenRouter	OpenRouter leads overall (3/5 metrics)

12. Limitations, Missing Data, and Next Experiments

Limitation	Impact	Next step	Current handling
Full routing trace	Cannot prove every provider choice, fallback, and retry path hop by hop	Add provider routing trace, decision logs, and fallback reasons	Use only returned provider fields and provider distribution
Longer time window	4x50 observes short-window stability but not day-level drift	Add soak tests and repeated windows	Scope conclusions to this run
Production corpus	Built-in templates do not cover every workload distribution	Use sanitized production-stratified corpora	Discuss representative long-context templates only
Cost-field consistency	Cost coverage and semantics may differ by platform	Reconcile with billing and provider cost breakdown	Use only explicitly returned cost fields

13. Reproducibility Appendix

Artifact	Path
Summary	export/qwen3_7_flash_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions summary.json
Paired dataset	export/qwen3_7_flash_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions benchmark_pairs.csv
Request- level dataset	export/qwen3_7_flash_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions benchmark_requests.jsonl
Filtered structured records	export/qwen3_7_flash_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions records.json
Excluded- record audit	export/qwen3_7_flash_all_experiments/ routing_sort_cache_cost_ab_4x50_stream_ttft_reasoning_default_length_tiers_chat_completions records_excluded.json
Test source	tests/test_rerun_routing_sort_cache_cost_ab.py
Benchmark runner source	scripts/rerun_routing_sort_cache_cost_ab.py
HTML report renderer source	scripts/render_glm52_deepseek_style_report.py
Dataset reference	business_representative built-in representative business templates; request-level export is benchmark_requests